



E&ICT Academy, IIT Kanpur

A Joint Initiative of MeitY & IIT Kanpur

Strategic Partner



CTO's Financial Toolkit

Session 7

Mentor:

Shubhashish Biswas

Introduction



Shubhashish Biswas

20 years experience in :

- **Managing Director, Algo & Consulting Advisor**, BCG
- **Practice Head, Digital & AI strategy** at KPMG, Accenture Strategy
- **VP (AI and Adv. Analytics)** at Citi
- **Product Management for Microsoft** (Machine Learning Offering on Microsoft Azure (Cloud))



POST GRADUATION

- **IIM Ahmedabad**
MBA (2-Year Full-time)
Focus: Technology / Systems & Finance
- **MIT, Cambridge (US),**
Pro. Masters Program
(Focus: Analytics & Data Science)

GRADUATION

IIT Madras
Bachelors Degree, Dept. of CSE
(Focus: AI & Analytics)



IIM
Ahmedabad
(2 Yr. PGP)



Massachusetts
Institute of
Technology
Cambridge, US



IIT Madras
Computer Science Dept.*



Session 7 Agenda



7.1

CTO's financial toolkit – A Refresher

7.2

Tool Deep Dive & Mini Cases

7.3

Capstone Project - AMA

The CTO's financial toolkit ecosystem

CXO

Enterprise Transformation

Driving organizational change through technology

AI Integration

Scaling AI adoption for maximum value

Risk Management

Mitigating cybersecurity threats and financial exposures

Business Strategy

Aligning technology with overall business goals

Cloud & Data Infrastructure

Optimizing cloud and data for efficiency and sustainability

Leadership Alignment

Unifying engineering, security, and finance teams

The Evolving CTO Role in 2026

Ref: [CTO Mandate in 2026 – CXO India](#)

Eight questions for the chief transformation officer



1

Do I have the full backing and confidence of the CEO and the board?



5

Am I coaching the CEO and top team and successfully changing how they lead the transformation?



2

Have I taken on incumbent vested interests and slaughtered relevant sacred cows?



6

Do I have a clear-eyed view on where the real value lies and where we cannot afford to compromise?



3

Have I established a rhythm and clock speed that is changing the metabolic rate of the business?



7

Have I intentionally picked some fights with senior line leaders and gotten them to change?




4

Have I connected to the front line and truly understood its perspectives and frustrations?



8

Do I understand the dominant culture and mind-set and where it needs to shift?



The CTO's financial toolkit ecosystem



Planning & Governance

- Budget Forecasting
- Investment Portfolio
- Chargeback Model



Investment & Spend Decision-making

- TCO Calculation
- Business Case Dev.
- Vendor Scorecard




Operational Efficiency

- Cloud Spend (FinOps)
- Tech Debt Register
- Unit Economics



Risk Management

- Risk Register
- M&A Due Diligence



People

- Workforce Model



E&ICT Academy, IIT Kanpur
A Joint Initiative of MeitY & IIT Kanpur

Strategic Partner
leadup
u n i v e r s e

01

Budget Forecast



Background

PERSPECTIVE

In 2025–26, technology budgets face three simultaneous pressures: unplanned AI investment demands arriving mid-year, cloud spend swinging $\pm 30\%$ quarter-to-quarter, and CFO scrutiny that requires real-time justification of every dollar.

CTOs operating on a single annual budget submission lose credibility fast — they can't respond to board questions in real time, and they can't reallocate to high-priority initiatives without a protracted approval cycle.

A rolling 12–18 month forecast, refreshed monthly, changes the CTO's posture from reactive to strategic. It forces the team to own their numbers — increasingly a prerequisite for a seat at the executive table.

CORE ELEMENTS

1. Personnel & contractors (by team, role, FTE vs. contractor split)
2. Cloud & infrastructure (committed vs. on-demand breakdown, by product/environment)
3. Software licenses & SaaS (annualized, renewal dates flagged)
4. Security & compliance spend (explicit named category, not buried in IT overhead)
5. AI/ML platform costs (separate board-visible line item)
6. CapEx vs. OpEx classification (critical for EBITDA and accounting treatment)
7. Contingency reserve (10–15% of discretionary spend, pre-approved with a release process)



E&ICT Academy, IIT Kanpur

A Joint Initiative of MeitY & IIT Kanpur

Strategic Partner



Tool Highlights

Budget Category	Q1	Q2	Q3	Q4	Annual Total	vs. Prior Year
PERSONNEL						
Engineering FTEs	\$1,200K	\$1,200K	\$1,280K	\$1,280K	\$4,960K	8%
Contractors & Staff Aug.	\$180K	\$200K	\$200K	\$180K	\$760K	15%
Training & Certification	\$40K	\$20K	\$40K	\$20K	\$120K	New
INFRASTRUCTURE & CLOUD						
Cloud Committed (RIs/Savings Plans)	\$420K	\$420K	\$420K	\$420K	\$1,680K	5%
Cloud On-Demand / Variable	\$120K	\$140K	\$130K	\$110K	\$500K	22%
Data Center / Co-location	\$80K	\$80K	\$80K	\$80K	\$320K	0%
SOFTWARE & SAAS LICENSES						
Productivity & Collaboration	\$60K	\$60K	\$60K	\$60K	\$240K	2%
Developer Tools	\$45K	\$45K	\$45K	\$45K	\$180K	10%
Security & Monitoring Tools	\$90K	\$90K	\$90K	\$90K	\$360K	20%
AI / ML PLATFORMS						
Foundation Model APIs	\$80K	\$100K	\$120K	\$130K	\$430K	New
MLOps & Data Platform	\$60K	\$60K	\$60K	\$60K	\$240K	New
SECURITY & COMPLIANCE						
Penetration Testing	\$50K	\$0	\$50K	\$0	\$100K	25%
Compliance Tooling (SOC2 etc.)	\$30K	\$30K	\$30K	\$30K	\$120K	10%
CONTINGENCY (12%)	\$280K	\$290K	\$305K	\$300K	\$1,175K	
TOTAL OPEX	\$2,735K	\$2,735K	\$2,910K	\$2,905K	\$11,285K	11%
CAPEX (Hardware, Licenses >1yr)	\$400K	\$0	\$200K	\$0	\$600K	-20%
TOTAL BUDGET	\$3,135K	\$2,735K	\$3,110K	\$2,905K	\$11,885K	9%



Mini Case



TechCorp is a 3,000-person B2B SaaS company. The CTO is presenting FY2026 budget to the CFO and CEO. Midway through FY2025, the board asked for an AI strategy — which now needs funding. CFO's opening position is: 'I want 5% cost reduction in tech overall.'

What is your strategy to set the budget forecast for it?



Mini Case Solution Playbook

Establish forecast credibility first

Use the rolling forecast to show Q3/Q4 actuals vs. budget with accuracy within 5%. A CTO who forecasts accurately earns the benefit of the doubt on new investment requests before any new ask lands.

Frame the tradeoff with three scenarios

Present: (A) flat budget with no AI — maintains status quo, risks competitive lag; (B) 5% reduction with reallocation — cuts SaaS redundancy (\$400K savings) to fund AI (\$430K); (C) 9% increase — full AI platform plus 2 engineers. Recommend (B) explicitly.

Show AI as a named budget line

Use the template's AI section to show exactly what the spend is, what it produces (productivity per engineer, feature velocity), and what 6-month and 12-month success metrics look like.

Self-fund with SaaS rationalization

Run a vendor audit first and Identify \$400K+ in underutilized licenses. Use this to fund the AI line, reducing the net ask to near zero.

Outcome

Board approves Scenario B with a mid-year review trigger at \$200K of AI spend. CTO has a monthly dashboard tracking actuals vs. forecast with automated alerting at 80% of quarterly budget.



E&ICT Academy, IIT Kanpur

A Joint Initiative of MeitY & IIT Kanpur

Strategic Partner



02

Investment Portfolio Tracker



Tool 2: Investment Portfolio Tracker

Introduction

PERSPECTIVE

The Run/Grow/Transform framework divides technology investment into three buckets: keeping the lights on (Run), expanding existing capabilities (Grow), and creating the future (Transform). Most enterprises spend 70–80% on Run without realizing it, starving Grow and Transform.

CTOs who manage a technology investment portfolio rather than a project list make fundamentally better allocation decisions. They can defend 'why so much on infrastructure?' and 'why only \$200K on AI?' with data, not intuition.

In 2025, AI is the most contested item in Transform budgets. The portfolio lens helps CTOs allocate AI investment deliberately — not reactively — and show the board that technology strategy is coherent, not ad hoc.

CORE ELEMENTS

1. Run portfolio (infrastructure, licenses, support — keeping the business operational)
2. Grow portfolio (product enhancements, market expansion, process automation)
3. Transform portfolio (AI, platform re-architecture, new business models)
4. Current allocation vs. industry benchmark vs. target allocation (side-by-side)
5. Portfolio health metrics (ROI by initiative, run rate vs. budget, on-time completion)
6. Strategic alignment score (each initiative linked to a named board priority)
7. 18–24 month rebalancing plan (the path from current to target allocation)



Tool 2: Investment Portfolio Tracker

Tool Snapshot

Category	Current (\$K)	Current (%)	Benchmark	Target (%)	Target (\$K)	Action Required
RUN (Keep the lights on)						
Infrastructure & Cloud	\$2,800K	24%			\$2,400K	Reduce via FinOps (Tool 04)
Legacy System Maintenance	\$3,200K	27%			\$2,200K	Reduce via Tech Debt (Tool 05)
Support & Operations	\$1,400K	12%			\$1,200K	Automate to reduce
RUN SUBTOTAL	\$7,400K	62%	50–60% target	55%	\$6,490K	Free \$910K over 18 months
GROW (Expand existing capabilities)						
Product Feature Development	\$2,600K	22%			\$3,000K	Increase with Run savings
Process Automation	\$600K	5%			\$800K	Increase
Data & Analytics	\$400K	3%			\$600K	Increase
GROW SUBTOTAL	\$3,600K	30%	25–35% target	30%	\$3,540K	Hold / slight increase
TRANSFORM (Create the future)						
AI / ML Platform & Features	\$430K	4%			\$900K	Major increase — board priority
Architecture Modernization	\$320K	3%			\$600K	Increase as debt clears
New Business Model Experiments	\$50K	1%			\$270K	Create dedicated fund
TRANSFORM SUBTOTAL	\$800K	8%	10–20% target	15%	\$1,770K	Invest \$970K more/yr
TOTAL PORTFOLIO	\$11,800K	100%		100%	\$11,800K	Rebalance, not increase



Mini Case



ValNext Board Asks: 'Why Does Tech Cost So Much and Produce So Little?'

The board of a \$150M ARR SaaS company ValNext challenges the CTO: engineering headcount grew 30% last year but feature velocity feels flat. 'Where is all the money going?' The CTO has the board meeting in 10 days. What will you do in the 10 days?



Mini Case Solution Playbook

Map the portfolio in 48 hours

Use finance data and engineering team input to classify every initiative as Run, Grow, or Transform. It is often shocking: most CTOs discover they are at 65–75% Run without knowing it — the data makes the case for them.

Show the Run creep narrative

Present how the Run percentage increased from 55% two years ago to 62% today — driven by technical debt accumulation and cloud cost growth. This explains velocity decline without sounding defensive.

Present the rebalancing plan with specific tools

Show target allocation (55/30/15) and the 18-month roadmap: FinOps savings (\$1.1M from Tool 04), tech debt remediation (\$1.4M from Tool 05), vendor rationalization (\$400K from Tool 07). Total reallocation: \$2.9M toward Grow and Transform.

Map Transform investment to board objectives

Connect each Transform initiative to a board OKR: AI underwriting maps to loss ratio improvement; new API platform maps to partner ecosystem revenue. The board sees strategy, not a wish list.

Outcome

Board approves the rebalancing plan and adds quarterly technology portfolio review to the governance calendar. CTO now presents a technology P&L — not just headcount and spend — at every board meeting.



E&ICT Academy, IIT Kanpur

A Joint Initiative of MeitY & IIT Kanpur

Strategic Partner



03

IT Chargeback & Showback Model



Tool 3: Chargeback & Showback Models

PERSPECTIVE

When technology costs are pooled and invisible to business units, each team rationally maximizes its own consumption: spin up more resources, request more features, onboard more users — because the cost falls on IT's budget, not theirs. Chargeback makes each team accountable for what they consume.

Showback (reporting without billing) is a softer first step: teams see their monthly consumption costs, which alone drives 15–25% waste reduction even without a billing mechanism. For most organizations, showback is the right first step before full chargeback.

In 2025, AI cost governance is the newest and most urgent application: when an engineering team can run 10,000 foundation model API calls in an afternoon with no visible cost, they will. AI costs must be allocated by team and product — via showback at minimum — before they become uncontrollable.

CORE ELEMENTS

1. Service catalog (list of all IT services with a defined unit price for each)
2. Consumption tracking by business unit (cloud resources, licenses, support tickets)
3. AI usage tracking per team and per application (tokens, GPU hours, API requests)
4. Monthly allocation report (showback: 'your team consumed \$X this month')
5. Chargeback mechanism (how and when costs are transferred to business unit budgets)
6. Dispute resolution process (how business units challenge incorrect allocations)
7. Governance model (who sets service prices, how often reviewed, who adjudicates disputes)

Tool Snapshot

IT Service Catalog with Unit Pricing

USAGE NOTES

- Start with showback for 2–3 quarters before moving to formal chargeback — teams need time to build cost awareness.
- Set unit prices at 100–110% of actual cost to recover the administrative overhead of operating the allocation process.
- AI service costs are the fastest-changing line item — review and update pricing monthly, not quarterly.
- Never allocate shared infrastructure (email, VPN, identity) via chargeback — use a flat key (headcount %) for genuinely shared services.

Service	Unit	Unit Price	Primary Consumer	Jan Usage	Jan Allocation	Notes
INFRASTRUCTURE						
AWS Compute (EC2 equivalent)	vCPU-hour	\$0.06	Product, Data	42,000 hrs	\$2,310	Tagged by team in cloud console
Cloud Storage (S3 equivalent)	GB-month	\$0.03	Data, Analytics	180,000 GB	\$4,500	Uncompressed storage
Managed Database (RDS equivalent)	DB-instance-hour	\$0.22	Backend Eng.	4,400 hrs	\$968	Per active instance
DEVELOPER PLATFORM						
GitHub Enterprise (per seat)	seat-month	\$21	Engineering	180 seats	\$3,780	Annual license amortized monthly
CI/CD Pipeline Minutes	compute-minute	\$0.01	Engineering	180,000 min	\$1,440	Build job minutes
Monitoring (Datadog)	host-month	\$28	All teams	85 hosts	\$2,380	Per active monitored host
AI / ML SERVICES						
Foundation Model API (GPT-4 class)	1,000 tokens	\$0.03	Product, CX	4.2M tokens	\$105	Tagged by team + feature
GPU Compute (training jobs)	GPU-hour	\$2.80	ML Team	320 GPU-hrs	\$896	On-demand GPU pool
SUPPORT SERVICES						
Help Desk (L1/L2 support)	ticket	\$45	All BUs	280 tickets	\$12,600	Attributed to requestor's BU
Security Reviews (per feature)	review	\$800	Product, Legal	4 reviews	\$3,200	Per feature security review
SAMPLE MONTHLY SHOWBACK: Product Team						
Infrastructure (compute, DB, storage)					\$4,820	
Developer platform (seats, CI, monitoring)					\$2,100	
AI / ML services (API + GPU)					\$3,640	
Support services (tickets raised)					\$1,350	
PRODUCT TEAM TOTAL — JANUARY					\$11,910/month	



Mini Case



In a mid-market tech products firm, three product teams are independently using foundation model APIs with no central visibility. The monthly AI API bill jumped from \$28K to \$112K in 90 days. No team knows how much they are spending. The CTO needs governance without bureaucracy that slows product development.

How will you govern this as the CTO?



Tool 3: IT chargeback and showback models

Mini-case

Tag and attribute all AI costs — Week 1–2

Tag all AI API calls with team and feature identifiers. Most cloud providers and API gateways support this natively. Without attribution, you cannot have a productive cost conversation — you can only say 'we spent \$112K' not 'your team spent \$X.'

Publish the first showback report

Share monthly AI cost per team in a visible, non-punitive format. In most cases, teams are shocked by their own usage. This alone drives 20–30% voluntary reduction within one billing cycle, with no mandates required.

Set per-team budgets with automated alerting

Agree on monthly AI spend budgets with each team lead — negotiated, not imposed. Implement alerting at 70% and 90% of budget. Teams retain agency while accountability is established.

Track AI cost efficiency, not just cost

Measure 'AI cost per feature delivered' or 'AI cost per user interaction' monthly. This reframes the conversation from 'you are spending too much' to 'here is how to spend productively — let us optimize together.'

Outcome

90 days after showback launch, total AI spend drops from \$112K to \$74K/month with no features removed. Teams implement caching, model-tier selection, and streaming optimization. CTO presents this as 34% efficiency improvement — not a cost cut.

04

Total Cost of Ownership (TCO) Calculator



Tool 4: TCO Calculator

PERSPECTIVE

CTOs routinely face decisions that look cheap in Year 1 but expensive by Year 3: an open-source tool requiring 2 dedicated engineers to maintain, a SaaS platform whose per-seat pricing explodes with growth, or an on-prem investment that can't scale without another data center build.

In 2025, the build vs. buy calculus has shifted again: AI tools reduce engineering effort, but also create new maintenance surfaces. The TCO model must now account for AI-augmented development costs, not just traditional FTE ratios.

A rigorous 5-year TCO protects the CTO from the most common trap: optimizing for the lowest Year 1 cost while ignoring the compounding cost of lock-in, integration debt, and scaling penalties.

CORE ELEMENTS

1. Acquisition / implementation costs (Year 0, one-time)
2. Annual operating costs (licensing, hosting, vendor support)
3. Internal engineering cost (integration + ongoing maintenance FTEs)
4. Scaling cost model at 2x and 5x current volume
5. Migration / exit cost (what it costs to leave the vendor or deprecate the build)
6. Opportunity cost of engineering hours diverted from product work
7. Risk-adjusted adjustments (vendor stability, technology obsolescence probability)

Tool 4: TCO Calculator

TCO Calculator: Build vs. Buy — 5-Year View

USAGE NOTES


- Opportunity cost = engineering hours × \$80/hr blended rate diverted from product features that generate revenue.
- Buy wins here due to high maintenance burden of custom data platforms. The calculus reverses for core IP that differentiates your product.
- Always include exit/migration cost in the Buy column — typically 6–12 months of engineering time to switch vendors.
- Recalculate annually: vendor pricing models change significantly, especially for AI-adjacent and data tooling.

Cost Component	Build (Internal)	Buy (Databricks+Snowflake)	Verdict
YEAR 0: IMPLEMENTATION			
Engineering time (build)	\$480K	\$80K	Buy wins
Infrastructure setup	\$120K	\$20K	Buy wins
Training & certification	\$40K	\$60K	Build wins
Total Year 0	\$640K	\$160K	Buy saves \$480K
YEARS 1–3: ANNUAL OPERATING (avg/yr)			
Licensing / SaaS subscription	\$0	\$480K/yr	Build wins (on paper)
Cloud compute costs	\$360K/yr	Included in license	Comparable
Engineering maintenance (FTEs)	\$640K/yr (2 FTEs)	\$80K/yr (0.25 FTE)	Buy wins decisively
Vendor support & SLA	\$0 (internal)	\$60K/yr	Build wins
Annual Operating Total	\$1,000K/yr	\$740K/yr	Buy saves \$260K/yr
YEARS 4–5: SCALE (5x volume)			
Additional engineering headcount	+\$640K (1 FTE)	\$0 (scales in license)	Buy wins
Infrastructure / compute scaling	\$400K	\$240K	Buy wins
License scaling (usage-based)	0	+\$180K	Build wins
5-YEAR TOTAL COST OF OWNERSHIP			
Total Cost (Year 0–5)	\$6,640K	\$4,560K	Buy saves \$2.08M
Engineering hours consumed	~18,000 hours	~2,500 hours	Buy saves 15,500 hrs
Opportunity cost of eng. hours*	\$1,440K	\$200K	Buy saves \$1.24M
RISK-ADJUSTED 5-YEAR TCO	\$8,080K	\$4,760K	Buy wins by \$3.32M



ROI & Business Case Framework

Mini Case



Evaluating a Custom Data Platform vs. Databricks + Snowflake

A \$200M ARR company's data engineering lead wants to build a custom lakehouse. A vendor team is pitching Databricks + Snowflake.

You as the CTO have 6 weeks to make a call and present to the CEO and CFO.



ROI & Business Case Framework

Mini-case solution playbook

Gather real cost inputs, not vendor estimates

Use current engineering salary data (blended FTE cost including benefits: 1.4–1.6x base), current cloud invoices by service, and actual support ticket volume. Build your own numbers — not the vendor's comparison spreadsheet.

Add a 50% buffer on all build estimates

Build work consistently underestimates by 40–60%. Apply a 50% buffer on engineering estimates for the build option. Explicitly model what 2 engineers NOT working on product features costs in delayed revenue.

Get vendor pricing in writing before comparing

Vendors frequently offer 30–50% discounts in competitive situations. Negotiate before the TCO is presented to leadership. You want the best realistic Buy cost — not list price.

Lead with the opportunity cost number

'Those 18,000 engineering hours could ship 4 major product features or reduce time-to-market by 6 months.' This resonates with the CEO more than raw dollar comparison alone.

Outcome

CTO selects Buy with a 3-year committed deal at 35% discount, reducing Buy TCO to \$3.8M vs. Build's \$8M. Engineering team redeployed to the product roadmap, shipping 2 major features within the year.

05

ROI & Business Case Framework



Introduction

PERSPECTIVE

Technology initiatives fail to get funded not because they lack merit, but because they are presented in engineering language: story points, uptime percentages, and capability descriptions that mean little to a CFO. A structured business case translates technical investment into financial outcomes.

In 2025, two dynamics raise the stakes: AI investments require a business case because costs are visible (API spend, platform licenses) but benefits are diffuse; and higher cost of capital means payback periods matter more. An 18-month payback beats a 36-month payback even at higher NPV.

CTOs who master the business case framework earn a different kind of credibility — they become investment allocators, not just technology managers. That is the path to a real seat in strategic conversations.

CORE ELEMENTS

1. Executive summary (the problem, solution, the ask, and the payback period)
2. Investment required: one-time + ongoing costs, CapEx vs. OpEx clearly separated
3. Quantified benefits: revenue impact, cost savings, risk reduction — all in dollars
4. Payback period (months to break-even — the top metric for most boards in 2025)
5. Net Present Value (NPV) at the company's cost of capital / hurdle rate
6. Internal Rate of Return (IRR) — the annualized return on the investment
7. Sensitivity / stress test at 50% of forecast benefits
8. Go/No-Go decision criteria with named milestones and explicit stop triggers



Tool Snapshot

EXECUTIVE SUMMARY	
The Problem	Customer onboarding takes 14 days (industry avg: 3 days). We lose ~12% of new customers during onboarding. At \$1.2M new bookings/month, this is \$144K/month in churn.
The Solution	Automated onboarding platform: API integrations, guided workflows, ML-based at-risk detection, and real-time progress dashboards for customers.
The Ask	\$850K total investment in Year 1. Expected payback: 7 months (base case).
INVESTMENT REQUIRED	
Engineering (12 months, 3 FTEs)	\$480K
Third-party tooling & APIs	\$120K
Infrastructure (cloud)	\$60K
Project management & QA	\$80K
Contingency (15%)	\$110K
Total Investment	\$850K Year 1; ongoing cost: \$120K/yr
QUANTIFIED BENEFITS	
Onboarding churn reduction (60% of current 12% rate)	\$86K/month → \$1,030K/year
Support cost reduction (fewer onboarding tickets)	\$15K/month → \$180K/year
Engineering time freed (fewer manual escalations)	\$25K/month → \$300K/year
Total Annual Benefit (recurring)	\$1,510K/year
FINANCIAL METRICS	
Payback Period	7 months (base) / 13 months (50% benefit stress test)
3-Year NPV (10% hurdle rate)	\$2,870K
Internal Rate of Return (3-year)	174%
Stress Test: benefits at 50% of forecast	Still pays back in 13 months — investment is robust
DECISION CRITERIA	
Go Trigger	Board approval + engineering capacity confirmed for Q1 start
6-Month Success Metrics	Onboarding time under 5 days; onboarding churn below 6%
Stop Trigger	At Month 3, if onboarding time not reduced 30%, pause and reassess scope



E&ICT Academy, IIT Kanpur

A Joint Initiative of MeitY & IIT Kanpur

Strategic Partner



Tool x: NAME

Mini Case



Proposing a \$3M AI Implementation to a Skeptical CFO

A financial services CTO wants to implement AI-assisted underwriting across three product lines. The ask is \$3M over 18 months. The CFO has seen overpromised AI projects fail at two previous companies. The board has requested an AI strategy but hasn't committed budget.

How should the CTO make the case?



Tool x: NAME

Mini Case Solution Playbook

Start with a \$400K pilot, not the full ask

Present a 90-day pilot on one product line with a clear success metric. Fast, measurable wins build credibility. The full \$3M business case is far easier to approve after a pilot you can point to.

Anchor benefits to existing data

'Underwriters currently process 40 applications/day. The AI pilot targets 65/day — a 62% throughput improvement we will measure directly.' This is a testable hypothesis, not a promise.

Quantify the cost of not acting

At current throughput, \$2.4M in annual premiums sits on the table due to processing delays. Frame both ways: here is what we gain, and here is what we lose each month by not moving.

Use IRR to handle the time objection

When the CFO says '18 months is a long time,' IRR of 174% reframes it: 'We are earning 174% annually on this investment — 17x our cost of capital. The question is whether we can execute fast enough to capture the return.'

Outcome

Board approves the \$400K pilot with pre-authorization for the full \$3M if it hits 50% of projected throughput improvement. Pilot succeeds; full program approved at Month 4 — 14 months ahead of the original timeline.

06

Vendor Scorecard



Vendor Scorecard

PERSPECTIVE

The average enterprise runs 200–300 SaaS applications. Procurement decisions were made by different teams at different times, often without coordination. The result: overlapping capabilities, unused licenses, and auto-renewed contracts nobody remembers signing.

In 2025, the pressure is acute: CFOs are running SaaS rationalization audits in response to macroeconomic scrutiny, and the CTO is expected to lead the technical component. AI tooling (Zylo, Torii, Productiv) makes it possible to measure actual utilization vs. licensed seats in days.

Beyond cost savings, vendor scorecard discipline forces strategic clarity: if the company runs 5 different monitoring tools, that is not just waste — it is a signal of underdeveloped architectural standards. The vendor portfolio reveals strategy gaps hiding in plain sight.

CORE ELEMENTS

1. Complete vendor inventory (all contracts, annual spend, renewal dates, named tool owners)
2. Utilization rate (licensed seats vs. active users in last 90 days, by tool)
3. Strategic fit score (essential / useful / redundant / legacy)
4. Switching cost estimate (migration effort + productivity loss if consolidated)
5. Negotiation leverage assessment (competitive alternatives, market pricing)
6. Total SaaS spend as % of ARR (8–15% is typical for B2B SaaS)
7. Priority action matrix: renew / renegotiate / consolidate / eliminate

Tool Snapshot

Vendor / Tool	Annual Spend	Seats Licensed	Active Users (90d)	Utilization	Strategic Fit	Recommended Action
Salesforce (CRM)	\$480K	250	238	95%	Essential	Renew — negotiate 5% reduction
Slack	\$180K	800	720	90%	Essential	Renew at current terms
GitHub Enterprise	\$140K	200	185	93%	Essential	Renew at current terms
Datadog (monitoring)	\$210K	—	—	High	Essential	Renegotiate with New Relic competitive bid
PagerDuty	\$95K	150	48	32%	Redundant	Eliminate — consolidate into Datadog
Zendesk	\$120K	80	62	78%	Useful	Downgrade plan — save \$30K/yr
Confluence	\$90K	800	280	35%	Redundant	Review — Notion used by 400 users actively
Notion	\$75K	400	390	98%	Useful	Renew — migrate Confluence users here
Snowflake	\$320K	—	—	Medium	Essential	Optimize query costs — est. \$60K savings
Zoom	\$95K	800	750	94%	Essential	Renew at current terms
15 tools under \$50K each	\$420K	—	—	Mixed	Mixed	Audit individually — 30%+ elimination likely
TOTALS / SAVINGS SUMMARY	\$2,225K total					
Immediate savings (eliminate + downgrade)	-\$125K/yr					
Negotiation savings (competitive bids)	-\$85K/yr					
Query/usage optimization savings	-\$60K/yr					
TOTAL IDENTIFIED SAVINGS	-\$270K/yr (12% of SaaS spend)					

USAGE NOTES

- Utilization below 50% is a strong signal to consolidate or eliminate — below 30% is almost always waste.
- Bring competitive quotes to every renewal above \$100K — you do not need to switch, just demonstrate you could.
- SaaS spend benchmark: 8–15% of ARR for B2B SaaS. Above 15% warrants an immediate audit.
- Check for shadow IT: use expense report data and SSO logs to find unapproved tools teams are paying for personally.



Mini Case



\$2M in SaaS Budget Cuts Required by End of Quarter

The CFO announces company-wide cost reduction: every function must cut 15% of discretionary spend. Technology's target is \$2M over 12 months.

You, the CTO have 30 days to present a plan that does not cripple engineering velocity.

How will you plan it?



Mini Case Solution Playbook

Run the utilization audit immediately

Pull SSO login data, product analytics, and expense reports. Within 5 business days, identify every tool below 60% utilization. This is the no-regret zone — eliminating unused tools causes minimal engineering pain.

Segment by replaceability

Classify vendors as: (1) Essential with no substitute, (2) Essential with an available alternative, (3) Redundant — capability already covered by another tool, (4) Nice-to-have. Eliminations come from 3 and 4 first.

Negotiate all renewals proactively

Contact every vendor with a renewal in the next 6 months and tell them you are benchmarking alternatives. A 15–20% discount on a \$200K contract in exchange for a 2-year commitment is common in the current SaaS market.

Present with a vendor consolidation map

Show the CFO a before/after vendor landscape: fewer tools, consolidated capabilities, clear owners. This demonstrates strategic intent, not just cuts — and builds confidence that engineering will not be impaired.

Outcome

CTO identifies \$340K in no-pain eliminations, \$220K in negotiated discounts, and \$180K in plan downgrades (\$740K total from vendor rationalization). The remaining gap is closed via FinOps cloud optimization (Tool 04), bringing total to \$1.9M.

Thank You